

Notebook

February 11, 2025

[24]:

```
[53]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats
import statsmodels.api as sm
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.ensemble import RandomForestRegressor
from sklearn.svm import SVC
from sklearn.feature_selection import VarianceThreshold
import shap
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.metrics import accuracy_score, classification_report, \
    ↪confusion_matrix
from statsmodels.discrete.discrete_model import MNLogit
from numpy.linalg import LinAlgError
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, \
    ↪f1_score, confusion_matrix, roc_auc_score
from statsmodels.tsa.holtwinters import ExponentialSmoothing
from statsmodels.tsa.holtwinters import Holt
from sklearn.metrics import mean_squared_error
```

[]:

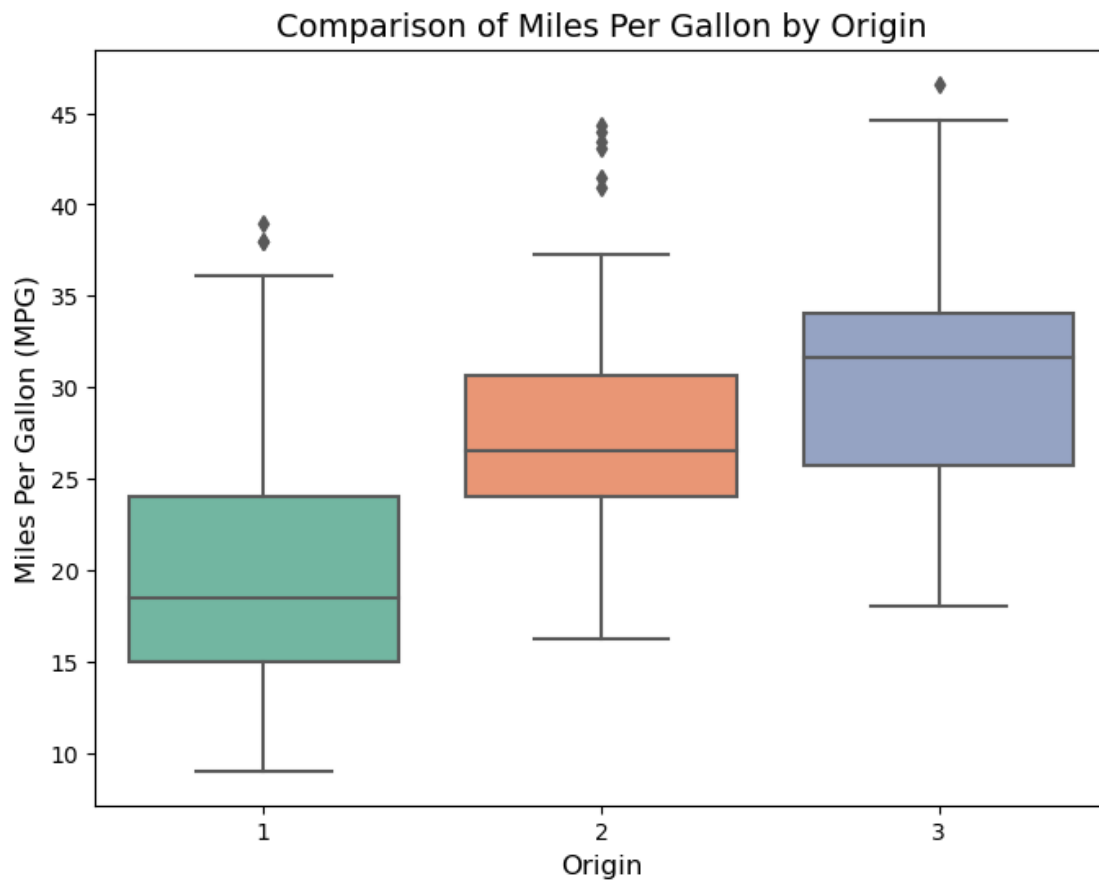
```
[54]: import os

# Get absolute path of the current notebook directory
BASE_DIR = os.path.join(os.getcwd(), "data")

# Function to get the full path
```

```
def get_file_path(filename):  
    return os.path.join(BASE_DIR, filename)
```

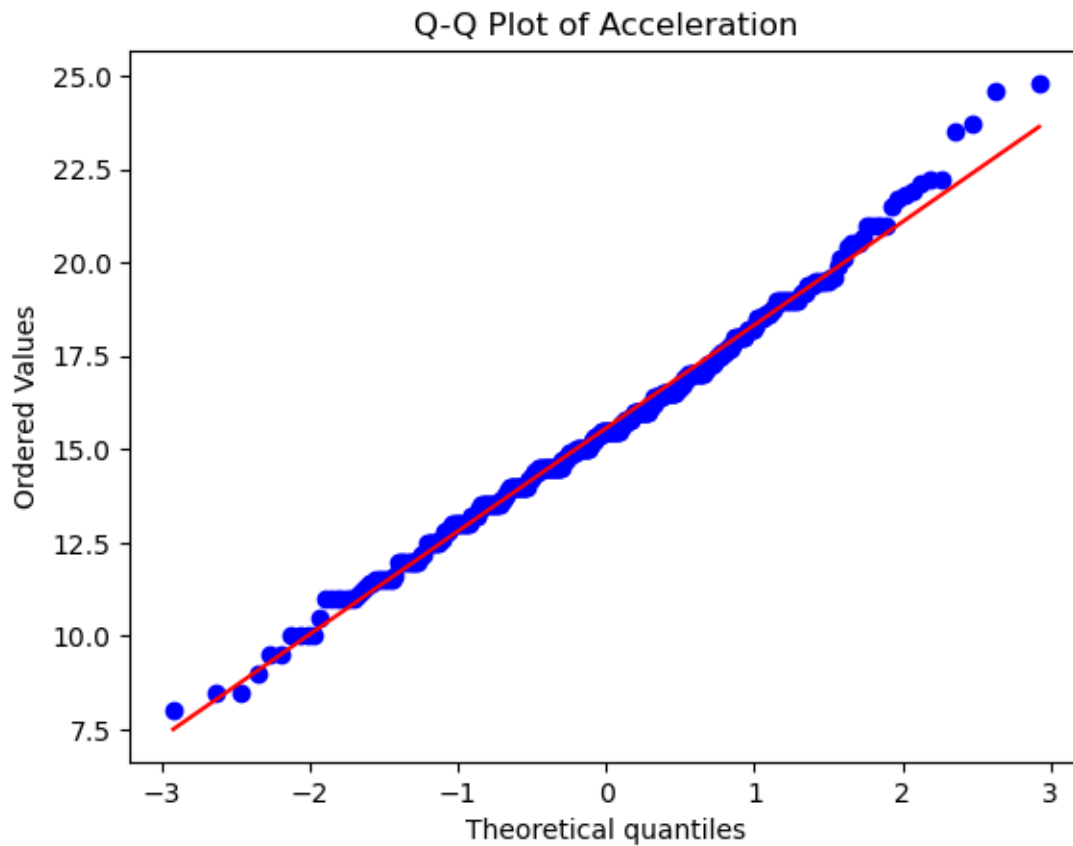
```
[55]: #load dataset  
df= pd.read_csv(get_file_path("2025_Box-Plot_Dataset.csv"))  
  
#creating a boxplot comparing Miles_Per_Gallon for different origin groups  
plt.figure(figsize=(8,6))  
sns.boxplot(x='Origin', y="Miles_Per_Gallon", data=df, palette="Set2")  
  
plt.title("Comparison of Miles Per Gallon by Origin", fontsize=14)  
plt.xlabel("Origin", fontsize=12)  
plt.ylabel("Miles Per Gallon (MPG)", fontsize=12)  
  
plt.show()
```



1 Q-Q Plot

```
[56]: # Q-Q plot to check if 'Acceleration' follows the normal distribution
```

```
stats.probplot(df['Acceleration'], dist="norm", plot=plt)
plt.title("Q-Q Plot of Acceleration")
plt.show()
```



1.1 Kolmogorov-Smirnov (K-S) Test

```
[57]: ks_stat, ks_p_value = stats.kstest(df['Acceleration'], 'norm',
    ↪ args=(df['Acceleration'].mean(), df['Acceleration'].std()))

print(f"Kolmogorov-Smirnov Test: Statistic={ks_stat:.4f}, p-value={ks_p_value:.
    ↪ 4f}")

alpha = 0.05 #level of significance

if ks_p_value <= alpha:
```

```

    print("K-S Test: Reject H0 → 'Acceleration' is NOT normally distributed.")
else:
    print("K-S Test: Fail to reject H0 → 'Acceleration' follows a normal_
↳distribution.")

```

Kolmogorov-Smirnov Test: Statistic=0.0508, p-value=0.2466

K-S Test: Fail to reject H0 → 'Acceleration' follows a normal distribution.

1.2 Shapiro-Wilk Test

```

[58]: shapiro_stat, shapiro_p_value = stats.shapiro(df['Acceleration'])

print(f"Shapiro-Wilk Test: Statistic={shapiro_stat:.4f},_
↳p-value={shapiro_p_value:.4f}")

alpha = 0.05 # Significance level

if shapiro_p_value < alpha:
    print("S-W Test: Reject H0 → 'Acceleration' is NOT normally distributed.")
else:
    print("S-W Test: Fail to reject H0 → 'Acceleration' follows a normal_
↳distribution.")

```

Shapiro-Wilk Test: Statistic=0.9924, p-value=0.0399

S-W Test: Reject H0 → 'Acceleration' is NOT normally distributed.

1.3 Linear Regression Model

[]:

```

[59]: df = pd.read_csv(get_file_path('2025_Regression_Dataset.csv'))

# Display first few rows
print(df.head())

```

	wtd_mean_atomic_radius	wtd_gmean_ThermalConductivity	\
0	105.514286	0.621979	
1	106.342857	0.624878	
2	104.371429	0.629441	
3	104.542857	0.633910	
4	104.885714	0.642942	

	wtd_gmean_FusionHeat	wtd_gmean_fie	wtd_gmean_ElectronAffinity	\
0	1.040986	938.016780	99.414682	
1	1.044545	937.025573	97.774719	
2	1.039211	940.294344	98.411962	
3	1.040986	940.391699	96.998357	
4	1.044545	940.586438	94.231770	

	wtd_entropy_ThermalConductivity	wtd_entropy_FusionHeat	\
0	0.262848	0.994998	
1	0.272820	1.044970	
2	0.283412	0.964031	
3	0.295609	0.994998	
4	0.316852	1.044970	

	wtd_entropy_ElectronAffinity	wtd_entropy_Density	std_FusionHeat	...	\
0	0.787382	0.814598	4.599064	...	
1	0.787396	0.859811	4.599064	...	
2	0.777657	0.769628	4.599064	...	
3	0.775688	0.791445	4.599064	...	
4	0.771173	0.830563	4.599064	...	

	mean_ThermalConductivity	mean_ElectronAffinity	mean_Density	\
0	107.756645	81.8375	4654.35725	
1	107.756645	81.8375	4654.35725	
2	112.006645	79.6075	4434.35725	
3	112.006645	79.6075	4434.35725	
4	112.006645	79.6075	4434.35725	

	gmean_ThermalConductivity	gmean_FusionHeat	gmean_fie	gmean_atomic_mass	\
0	7.062488	3.479475	718.152900	66.361592	
1	7.062488	3.479475	718.152900	66.361592	
2	8.339818	3.479475	734.219624	59.310096	
3	8.339818	3.479475	734.219624	59.310096	
4	8.339818	3.479475	734.219624	59.310096	

	entropy_ThermalConductivity	entropy_ElectronAffinity	critical_temp
0	0.308148	1.159687	29.0
1	0.308148	1.159687	23.0
2	0.403693	1.096672	36.0
3	0.403693	1.096672	31.0
4	0.403693	1.096672	33.0

[5 rows x 28 columns]

```
[60]: print(df.isnull().sum())

print(df.info())

print(df.describe())
```

wtd_mean_atomic_radius	0
wtd_gmean_ThermalConductivity	0
wtd_gmean_FusionHeat	0
wtd_gmean_fie	0

```

wtd_gmean_ElectronAffinity      0
wtd_entropy_ThermalConductivity  0
wtd_entropy_FusionHeat          0
wtd_entropy_ElectronAffinity     0
wtd_entropy_Density              0
std_FusionHeat                   0
std_fie                          0
std_atomic_radius                 0
std_atomic_mass                   0
range_Valence                     0
range_FusionHeat                  0
range_ElectronAffinity            0
range_Density                     0
mean_Valence                      0
mean_ThermalConductivity          0
mean_ElectronAffinity             0
mean_Density                      0
gmean_ThermalConductivity         0
gmean_FusionHeat                  0
gmean_fie                         0
gmean_atomic_mass                 0
entropy_ThermalConductivity        0
entropy_ElectronAffinity           0
critical_temp                     0
dtype: int64
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5275 entries, 0 to 5274
Data columns (total 28 columns):

```

#	Column	Non-Null Count	Dtype
0	wtd_mean_atomic_radius	5275 non-null	float64
1	wtd_gmean_ThermalConductivity	5275 non-null	float64
2	wtd_gmean_FusionHeat	5275 non-null	float64
3	wtd_gmean_fie	5275 non-null	float64
4	wtd_gmean_ElectronAffinity	5275 non-null	float64
5	wtd_entropy_ThermalConductivity	5275 non-null	float64
6	wtd_entropy_FusionHeat	5275 non-null	float64
7	wtd_entropy_ElectronAffinity	5275 non-null	float64
8	wtd_entropy_Density	5275 non-null	float64
9	std_FusionHeat	5275 non-null	float64
10	std_fie	5275 non-null	float64
11	std_atomic_radius	5275 non-null	float64
12	std_atomic_mass	5275 non-null	float64
13	range_Valence	5275 non-null	int64
14	range_FusionHeat	5275 non-null	float64
15	range_ElectronAffinity	5275 non-null	float64
16	range_Density	5275 non-null	float64
17	mean_Valence	5275 non-null	float64

18	mean_ThermalConductivity	5275	non-null	float64
19	mean_ElectronAffinity	5275	non-null	float64
20	mean_Density	5275	non-null	float64
21	gmean_ThermalConductivity	5275	non-null	float64
22	gmean_FusionHeat	5275	non-null	float64
23	gmean_fie	5275	non-null	float64
24	gmean_atomic_mass	5275	non-null	float64
25	entropy_ThermalConductivity	5275	non-null	float64
26	entropy_ElectronAffinity	5275	non-null	float64
27	critical_temp	5275	non-null	float64

dtypes: float64(27), int64(1)

memory usage: 1.1 MB

None

	wtd_mean_atomic_radius	wtd_gmean_ThermalConductivity \
count	5275.000000	5275.000000
mean	134.681841	26.831423
std	28.787768	40.105763
min	64.600000	0.072768
25%	112.140667	1.084912
50%	125.833333	6.085087
75%	158.391200	45.875927
max	253.000000	358.713959

	wtd_gmean_FusionHeat	wtd_gmean_fie	wtd_gmean_ElectronAffinity \
count	5275.000000	5275.000000	5275.000000
mean	9.990091	833.516080	72.219307
std	12.973263	119.359462	31.490091
min	0.480799	502.500000	1.500000
25%	1.321875	721.053648	50.438290
50%	4.821853	858.874330	72.854039
75%	16.393145	937.697803	89.962916
max	105.000000	1183.712294	214.651659

	wtd_entropy_ThermalConductivity	wtd_entropy_FusionHeat \
count	5275.000000	5275.000000
mean	0.542039	0.922645
std	0.319675	0.367980
min	0.000000	0.000000
25%	0.250750	0.679053
50%	0.551852	1.000424
75%	0.777388	1.163470
max	1.584219	1.674166

	wtd_entropy_ElectronAffinity	wtd_entropy_Density	std_FusionHeat ... \
count	5275.000000	5275.000000	5275.000000 ...
mean	0.777248	0.862335	8.271298 ...
std	0.285434	0.318776	8.666517 ...
min	0.000000	0.000000	0.000000 ...

25%	0.666039	0.690637	4.261726	...
50%	0.785771	0.892008	4.948155	...
75%	0.879506	1.089535	8.935301	...
max	1.675375	1.659095	51.635000	...

	mean_ThermalConductivity	mean_ElectronAffinity	mean_Density	\
count	5275.000000	5275.000000	5275.000000	
mean	89.798759	76.945910	6117.305246	
std	38.398988	27.810696	2866.393179	
min	0.115415	1.500000	535.000000	
25%	61.000000	61.713333	4529.571500	
50%	97.000000	73.100000	5329.085800	
75%	111.004430	85.512500	6642.000000	
max	332.500000	228.300000	22590.000000	


	gmean_ThermalConductivity	gmean_FusionHeat	gmean_fie	\
count	5275.000000	5275.000000	5275.000000	
mean	29.281409	10.020024	738.018816	
std	33.079523	9.744749	77.991794	
min	0.072768	0.640000	490.633468	
25%	8.339818	4.141514	692.541331	
50%	14.287643	5.196603	728.828771	
75%	41.279535	13.481005	765.779664	
max	317.883627	105.000000	1225.900159	

	gmean_atomic_mass	entropy_ThermalConductivity	\
count	5275.000000	5275.000000	
mean	71.218091	0.729671	
std	31.279883	0.325210	
min	5.320573	0.000000	
25%	58.041225	0.457810	
50%	66.361592	0.741854	
75%	77.118550	0.962398	
max	208.980400	1.633977	

	entropy_ElectronAffinity	critical_temp
count	5275.000000	5275.000000
mean	1.077926	35.006786
std	0.337212	34.121958
min	0.000000	0.000500
25%	0.894571	5.700000
50%	1.158144	21.500000
75%	1.349501	64.350000
max	1.767732	136.000000

[8 rows x 28 columns]


```
[61]: # Featuyre selection => Perform correlation analysis to check:
plt.figure(figsize=(10, 8))
sns.heatmap(df.corr(), annot=False, cmap='coolwarm')
plt.show()
```



```
[62]: #Selected indepdet variables based on corrrrelation
selected_features_1
↳=["wtd_gmean_fie","wtd_entropy_FusionHeat","wtd_entropy_ElectronAffinity","wtd_entropy_Dens
↳="std_fie","std_atomic_radius", "std_atomic_mass",
"range_ElectronAffinity","range_Density","mean_ThermalConductivity","entropy_ElectronAffinity"

# Define independent (X) and dependent (y) variables
X = df[selected_features_1]
y = df["critical_temp"]
```

```

#split data into training & testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)

# Initialize and train model
model = LinearRegression()
model.fit(X_train, y_train)

```

[62]: LinearRegression()

```

[63]: print(f"Intercept: {model.intercept_}")
coefficients = pd.DataFrame({"Feature": selected_features, "Coefficient": model.
↳coef_})
print(coefficients)

```

```

Intercept: 9.858758417343662

```

	Feature	Coefficient
0	wtd_gmean_fie	-0.050728
1	wtd_entropy_FusionHeat	45.883388
2	wtd_entropy_ElectronAffinity	-29.566577
3	wtd_entropy_Density	-0.058841
4	std_fie	0.091861
5	std_atomic_radius	0.413353
6	std_atomic_mass	0.126460
7	range_ElectronAffinity	-0.032388
8	range_Density	-0.000680
9	mean_ThermalConductivity	0.179015
10	entropy_ElectronAffinity	-4.652724

```

[64]: y_pred = model.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error: {mse}")
print(f"R-squared Score: {r2}")

```

```

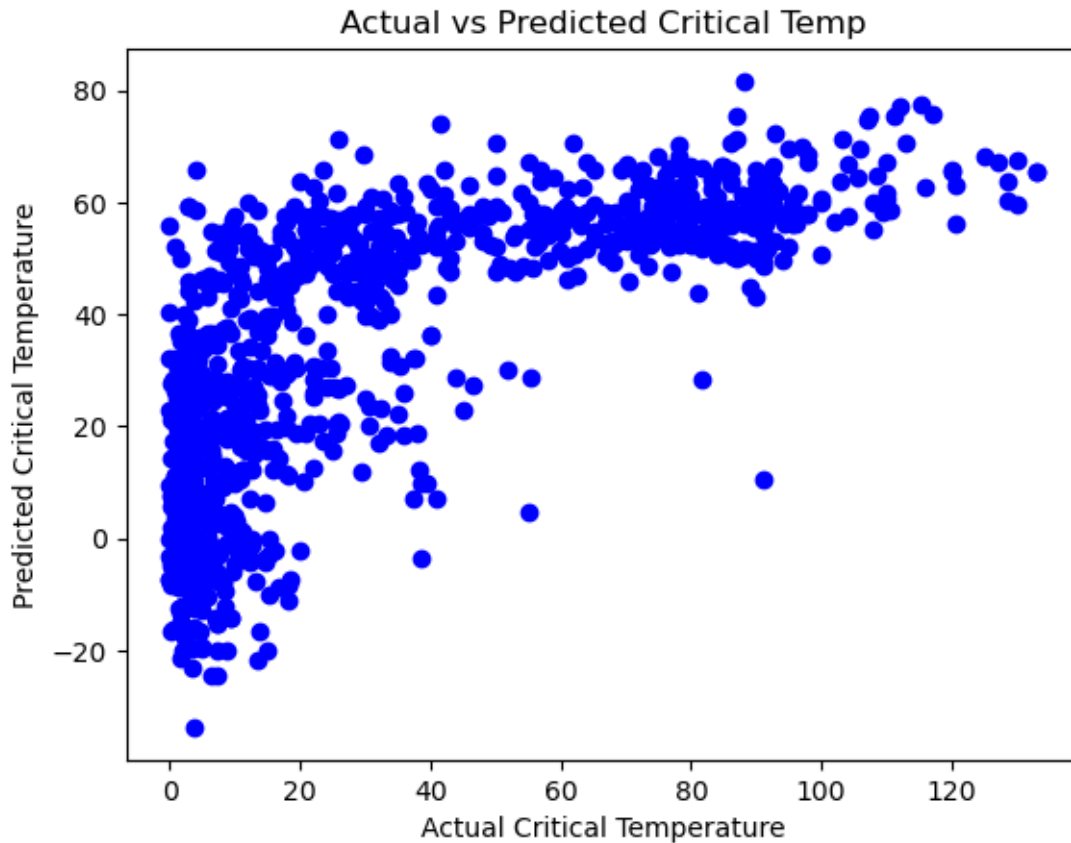
Mean Squared Error: 569.6824682665567
R-squared Score: 0.5159699869885219

```

```

[65]: plt.scatter(y_test, y_pred, color='blue')
plt.xlabel("Actual Critical Temperature")
plt.ylabel("Predicted Critical Temperature")
plt.title("Actual vs Predicted Critical Temp")
plt.show()

```



1.4 Logistic Regression Backward

```
[66]: df = pd.read_csv(get_file_path("2025_Classification_Dataset.csv"))

# Data preprocessing
X = df.drop(columns=['class'])
y = df['class']

# Split the dataset first to prevent data leakage
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
↳random_state=42)

# Handle categorical variables using LabelEncoder on training data and
↳transform test data
for col in X_train.select_dtypes(include=['object']).columns:
    le = LabelEncoder()
    # Fit on training data and transform both train and test
    X_train[col] = le.fit_transform(X_train[col].astype(str))
    # Handle unseen categories in test by assigning a default value (e.g., -1)
```

```

    # Note: This may require more sophisticated handling depending on the
    ↪ dataset
    X_test[col] = le.transform(X_test[col].astype(str))

# Remove constant features from training data
selector = VarianceThreshold(threshold=0)
X_train_processed = selector.fit_transform(X_train)
selected_columns = X_train.columns[selector.get_support()]
X_train = pd.DataFrame(X_train_processed, columns=selected_columns,
    ↪ index=X_train.index)
X_test = X_test[selected_columns]

# Add constant for statsmodels
X_train_sm = sm.add_constant(X_train)

# Backward Selection with intercept protection
def backward_elimination(X, y):
    X_opt = X.copy()
    while True:
        try:
            model = MNLogit(y, X_opt).fit(dis=0)
        except LinAlgError:
            # Handle singular matrix by removing the last added feature and
            ↪ break
            break
        p_values = model.pvalues
        # Exclude 'const' from consideration
        p_values_filtered = p_values.drop('const', errors='ignore')
        if p_values_filtered.empty:
            break
        max_p_value = p_values_filtered.max().max()
        if max_p_value > 0.05:
            # Find the feature with the highest p-value across any class
            max_feature = p_values_filtered.max(axis=1).idxmax()
            X_opt = X_opt.drop(columns=[max_feature])
        else:
            break
    return X_opt.columns

selected_features = backward_elimination(X_train_sm, y_train)
# Ensure 'const' is included if present
if 'const' in X_train_sm.columns:
    selected_features = list(selected_features) + ['const']
X_train_selected = X_train_sm[selected_features]
X_test_selected = sm.add_constant(X_test, has_constant='add')[selected_features]

# Standardize features

```

```

scaler = StandardScaler()
X_train_selected_scaled = scaler.fit_transform(X_train_selected)
X_test_selected_scaled = scaler.transform(X_test_selected)

# Model Fitting with increased max_iter and solver
model = LogisticRegression(max_iter=1000, solver='lbfgs',
    ↪multi_class='multinomial')
model.fit(X_train_selected_scaled, y_train)

# Model Evaluation
y_pred = model.predict(X_test_selected_scaled)
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))

```

/home/musiliandrew/anaconda3/lib/python3.10/site-packages/statsmodels/base/model.py:604: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
 warnings.warn("Maximum Likelihood optimization failed to "

Accuracy: 0.9415584415584416

Confusion Matrix:

```

[[724   3]
 [ 42   1]]

```

Classification Report:

	precision	recall	f1-score	support
No	0.95	1.00	0.97	727
Yes	0.25	0.02	0.04	43
accuracy			0.94	770
macro avg	0.60	0.51	0.51	770
weighted avg	0.91	0.94	0.92	770

```

[67]: import numpy as np

# Get coefficients and intercept
intercept = model.intercept_
coefficients = model.coef_[0] # model.coef_ returns a 2D array, we take the
    ↪first row

# Print results
print("Intercept:", intercept)
for feature, coef in zip(selected_features, coefficients):
    print(f"Feature: {feature}, Coefficient: {coef}, Exp(Coeff): {np.
    ↪exp(coef)}")

```

Intercept: [-1.62724828]

```

Feature: const, Coefficient: 0.0, Exp(Coeff): 1.0
Feature: shift, Coefficient: 0.24586199666849642, Exp(Coeff): 1.2787230933176406
Feature: gpuls, Coefficient: 0.167953831801793, Exp(Coeff): 1.182881997856601
Feature: nbumps2, Coefficient: 0.16101110715643788, Exp(Coeff):
1.1746980162960454
Feature: nbumps3, Coefficient: 0.17179863722619895, Exp(Coeff):
1.1874387031877787
Feature: const, Coefficient: 0.0, Exp(Coeff): 1.0

```

[]:

```

[68]: df = pd.read_csv(get_file_path('2025_Credit-card-clients.csv'))

# Encode categorical variables
categorical_cols = df.select_dtypes(include=["object"]).columns
label_encoders = {}
for col in categorical_cols:
    label_encoders[col] = LabelEncoder()
    df[col] = label_encoders[col].fit_transform(df[col])

# Split features and target
X = df.drop(columns=["Y"]) # Replace 'target' with actual target column name
y = df["Y"]

# Scale numerical features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
↳ random_state=42)

```

1.5 Model 1: Support Vector Machine (SVM)

```

[ ]: svm_model = SVC(kernel="rbf", probability=True)
svm_model.fit(X_train, y_train)

# Predictions
y_pred_svm = svm_model.predict(X_test)
y_prob_svm = svm_model.predict_proba(X_test)[: , 1]

# Evaluation
print(" SVM Results ")
print("Accuracy:", accuracy_score(y_test, y_pred_svm))
print("Precision:", precision_score(y_test, y_pred_svm))
print("Recall:", recall_score(y_test, y_pred_svm))
print("F1 Score:", f1_score(y_test, y_pred_svm))

```

```

print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_svm))
print("ROC-AUC Score:", roc_auc_score(y_test, y_prob_svm))
print("=" * 50)

```

1.6 Model 2: CHAID Decision Tree

```

[22]: decision_tree = DecisionTreeClassifier(criterion="entropy", max_depth=5,
      ↪min_samples_split=10) # CHAID Approximation
decision_tree.fit(X_train, y_train)

# Predictions
y_pred_tree = decision_tree.predict(X_test)
y_prob_tree = decision_tree.predict_proba(X_test)[: , 1]

# Evaluation
print(" Decision Tree (CHAID) Results ")
print("Accuracy:", accuracy_score(y_test, y_pred_tree))
print("Precision:", precision_score(y_test, y_pred_tree))
print("Recall:", recall_score(y_test, y_pred_tree))
print("F1 Score:", f1_score(y_test, y_pred_tree))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_tree))
print("ROC-AUC Score:", roc_auc_score(y_test, y_prob_tree))
print("=" * 50)

```

```

Decision Tree (CHAID) Results
Accuracy: 0.806
Precision: 0.6790123456790124
Recall: 0.3473684210526316
F1 Score: 0.4596100278551532
Confusion Matrix:
[[1447   78]
 [ 310  165]]
ROC-AUC Score: 0.7404210526315789
=====

```

```

[23]: ## Model 3: K-Nearest Neighbors (KNN)

knn_model = KNeighborsClassifier(n_neighbors=5, metric="euclidean")
knn_model.fit(X_train, y_train)

# Predictions
y_pred_knn = knn_model.predict(X_test)
y_prob_knn = knn_model.predict_proba(X_test)[: , 1]

# Evaluation
print(" KNN Results ")
print("Accuracy:", accuracy_score(y_test, y_pred_knn))

```

```

print("Precision:", precision_score(y_test, y_pred_knn))
print("Recall:", recall_score(y_test, y_pred_knn))
print("F1 Score:", f1_score(y_test, y_pred_knn))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_knn))
print("ROC-AUC Score:", roc_auc_score(y_test, y_prob_knn))
print("=" * 50)

```

Exception ignored on calling ctypes callback function: <function _ThreadPoolInfo._find_modules_with_dl_iterate_phdr.<locals>.match_module_callback at 0x7feb9d97f9a0>

Traceback (most recent call last):

```

File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 400, in match_module_callback
    self._make_module_from_path(filepath)
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 515, in _make_module_from_path
    module = module_class(filepath, prefix, user_api, internal_api)
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 606, in __init__
    self.version = self.get_version()
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 646, in get_version
    config = get_config().split()

```

AttributeError: 'NoneType' object has no attribute 'split'

Exception ignored on calling ctypes callback function: <function _ThreadPoolInfo._find_modules_with_dl_iterate_phdr.<locals>.match_module_callback at 0x7feb9d97f9a0>

Traceback (most recent call last):

```

File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 400, in match_module_callback
    self._make_module_from_path(filepath)
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 515, in _make_module_from_path
    module = module_class(filepath, prefix, user_api, internal_api)
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 606, in __init__
    self.version = self.get_version()
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 646, in get_version
    config = get_config().split()

```

AttributeError: 'NoneType' object has no attribute 'split'

KNN Results

Accuracy: 0.769

Precision: 0.5239852398523985

Recall: 0.29894736842105263

F1 Score: 0.3806970509383378

Confusion Matrix:


```
[[1396 129]
 [ 333 142]]
ROC-AUC Score: 0.6747113028472821
=====
```

```
[28]: import shap
from sklearn.ensemble import RandomForestClassifier

# 1. Train a Random Forest CLASSIFIER (not Regressor)
rf_model = RandomForestClassifier(
    n_estimators=100,
    class_weight="balanced", # Handle class imbalance
    random_state=42
)
rf_model.fit(X_train, y_train)

# 2. Use TreeExplainer (optimized for tree models)
explainer = shap.TreeExplainer(rf_model)
shap_values = explainer.shap_values(X_test)

# 3. Rank features by importance (using SHAP values for class 1)
shap_importance = pd.DataFrame({
    "Feature": X.columns,
    "SHAP Importance": np.abs(shap_values[1]).mean(axis=0)
}).sort_values("SHAP Importance", ascending=False)

print("Top Features by SHAP Importance:\n", shap_importance.head(10))

# 4. Plot SHAP summary (for class 1)
shap.summary_plot(shap_values[1], X_test, plot_type="bar")
```

```
-----
ValueError                                Traceback (most recent call last)
Cell In[28], line 17
    14 shap_values = explainer.shap_values(X_test)
    16 # 3. Rank features by importance (using SHAP values for class 1)
--> 17 shap_importance = pd.DataFrame({
    18     "Feature": X.columns,
    19     "SHAP Importance": np.abs(shap_values[1]).mean(axis=0)
    20 }).sort_values("SHAP Importance", ascending=False)
    22 print("Top Features by SHAP Importance:\n", shap_importance.head(10))
    24 # 4. Plot SHAP summary (for class 1)

File ~/anaconda3/lib/python3.10/site-packages/pandas/core/frame.py:778, in DataFrame.__init__(self, data, index, columns, dtype, copy)
    772     mgr = self._init_mgr(
    773         data, axes={"index": index, "columns": columns}, dtype=dtype,
    774         copy=copy
```

```

774     )
775 elif isinstance(data, dict):
776     # GH#38939 de facto copy defaults to False only in non-dict cases
--> 777     mgr = _
    dict_to_mgr(data, index, columns, dtype=dtype, copy=copy, typ=manager)
779 elif isinstance(data, ma.MaskedArray):
780     from numpy.ma import mrecords

File ~/anaconda3/lib/python3.10/site-packages/pandas/core/interals/constructio.
py:503, in dict_to_mgr(data, index, columns, dtype, typ, copy)
499     else:
500         # dtype check to exclude e.g. range objects, scalars
501         arrays = [x.copy() if hasattr(x, "dtype") else x for x in array.]
--> 503 return _
    arrays_to_mgr(arrays, columns, index, dtype=dtype, typ=typ, consolidate=copy)

File ~/anaconda3/lib/python3.10/site-packages/pandas/core/interals/constructio.
py:114, in arrays_to_mgr(arrays, columns, index, dtype, verify_integrity, typ,
consolidate)
111 if verify_integrity:
112     # figure out the index, if necessary
113     if index is None:
--> 114         index = _extract_index(arrays)
115     else:
116         index = ensure_index(index)

File ~/anaconda3/lib/python3.10/site-packages/pandas/core/interals/constructio.
py:677, in _extract_index(data)
675 lengths = list(set(raw_lengths))
676 if len(lengths) > 1:
--> 677     raise ValueError("All arrays must be of the same length")
679 if have_dicts:
680     raise ValueError(
681         "Mixing dicts with non-Series may lead to ambiguous ordering."
682     )

ValueError: All arrays must be of the same length

```

```

[34]: import pandas as pd

# Load dataset with correct delimiter and handle inconsistent formatting
df = pd.read_csv(get_file_path("TimeSeries_Dataset.csv"),
                 header=None,
                 names=["Year", "Month", "Sales"],
                 skipinitialspace=True) # Removes extra spaces

# Forward fill missing years

```

```

df["Year"].fillna(method="ffill", inplace=True)

# Clean and convert 'Sales' column
df["Sales"] = df["Sales"].astype(str).str.replace(",", "").str.strip(' ').
    ↪astype(int)

# Create a 'Date' column
df["Date"] = pd.to_datetime(df["Year"].astype(int).astype(str) + "-" +
    ↪df["Month"].astype(int).astype(str), format="%Y-%m")

# Set 'Date' as index and drop unnecessary columns
df.set_index("Date", inplace=True)
df.drop(columns=["Year", "Month"], inplace=True)

# Plot sales over time
plt.figure(figsize=(12, 6))
plt.plot(df.index, df["Sales"], label="Sales", color="blue")
plt.title("Sales Trend Over Time")
plt.xlabel("Year")
plt.ylabel("Number of Sales")
plt.legend()
plt.show()

# Decomposition to analyze trend and seasonality
decomposition = sm.tsa.seasonal_decompose(df["Sales"], model='additive',
    ↪period=12)
decomposition.plot()
plt.show()

```

/tmp/ipykernel_12841/500565962.py:10: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

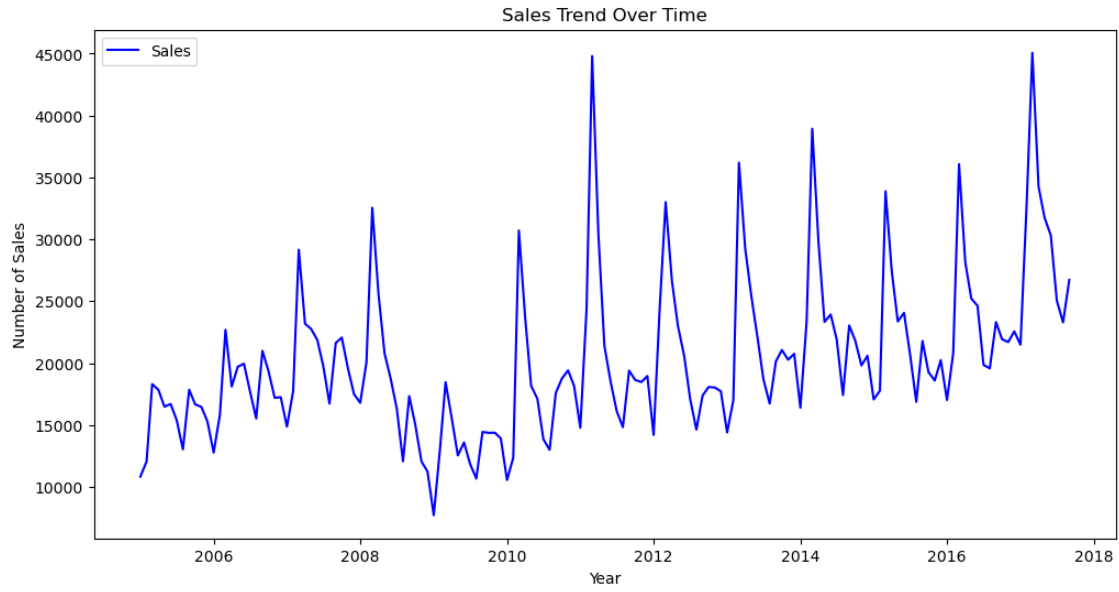
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```

df["Year"].fillna(method="ffill", inplace=True)
/tmp/ipykernel_12841/500565962.py:10: FutureWarning: Series.fillna with 'method'
is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill()
instead.
df["Year"].fillna(method="ffill", inplace=True)

```



```
[38]: # Splitting data
train = df[:'2015-12-01']
test = df['2016-01-01':]

# Holt's Linear Trend Model (HES)
holt_model = Holt(train['Sales']).fit()
holt_forecast = holt_model.forecast(len(test))
holt_mse = mean_squared_error(test['Sales'], holt_forecast)

# Holt-Winters Model (Multiplicative)
hw_model = ExponentialSmoothing(train['Sales'], trend='add', seasonal='add',
    ↪seasonal_periods=12).fit()
hw_forecast = hw_model.forecast(len(test))
hw_mse = mean_squared_error(test['Sales'], hw_forecast)

# Plotting
plt.figure(figsize=(10,5))
plt.plot(train.index, train['Sales'], label='Training Data')
plt.plot(test.index, test['Sales'], label='Actual Sales', color='black')
plt.plot(test.index, holt_forecast, label='Holt Forecast', linestyle='dashed')
plt.plot(test.index, hw_forecast, label='Holt-Winters Forecast',
    ↪linestyle='dotted')
plt.legend()
plt.show()

# Selecting the best model
best_model = "Holt-Winters" if hw_mse < holt_mse else "Holt"
print(f"MSE - Holt's Model: {holt_mse:.2f}")
print(f"MSE - Holt-Winters Model: {hw_mse:.2f}")
print(f"Recommended Model: {best_model}")
```

```
/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: No frequency
information was provided, so inferred frequency MS will be used.
```

```
self._init_dates(dates, freq)
```

```
/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/statsmodels/tsa/holtwinters/model.py:915: ConvergenceWarning:
Optimization failed to converge. Check mle_retvals.
```


```
warnings.warn(
```

```
/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: No frequency
information was provided, so inferred frequency MS will be used.
```

```
self._init_dates(dates, freq)
```

```
/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/statsmodels/tsa/holtwinters/model.py:915: ConvergenceWarning:
Optimization failed to converge. Check mle_retvals.
```

```
warnings.warn(
```



MSE - Holt's Model: 71802642.22
MSE - Holt-Winters Model: 16439804.83
Recommended Model: Holt-Winters


```
[40]: import scipy.cluster.hierarchy as sch
df = pd.read_csv(get_file_path("ClusterAnalysis_2025.csv"))

X = df.iloc[:, :].values

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

linked = sch.linkage(X_scaled, method='ward')

# Plot the dendrogram
plt.figure(figsize=(10, 5))
dendrogram = sch.dendrogram(linked)
plt.title('Dendrogram using Ward's method')
plt.xlabel('Observations')
plt.ylabel('Euclidean distances')
plt.show()
```



```
[43]: from sklearn.cluster import KMeans

k = 2

kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
df['Cluster'] = kmeans.fit_predict(X_scaled)

print("Cluster Centers:\n", kmeans.cluster_centers_)

plt.scatter(X_scaled[:, 0], X_scaled[:, 1], c=df['Cluster'], cmap='viridis',
            alpha=0.7)
plt.scatter(kmeans.cluster_centers_[:, 0], kmeans.cluster_centers_[:, 1],
            s=200, c='red', marker='X') # Cluster centers
plt.title('K-Means Clustering')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.show()
```

Exception ignored on calling ctypes callback function: <function _ThreadPoolInfo._find_modules_with_dl_iterate_phdr.<locals>.match_module_callback at 0x7feb7770cca0>

Traceback (most recent call last):

```
File "/home/musiliandrew/anaconda3/lib/python3.10/site-packages/threadpoolctl.py", line 400, in match_module_callback
    self._make_module_from_path(filepath)
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
```

```

packages/threadpoolctl.py", line 515, in _make_module_from_path
    module = module_class(filepath, prefix, user_api, internal_api)
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 606, in __init__
    self.version = self.get_version()
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 646, in get_version
    config = get_config().split()
AttributeError: 'NoneType' object has no attribute 'split'
Exception ignored on calling ctypes callback function: <function _ThreadPoolInfo
._find_modules_with_dl_iterate_phdr.<locals>.match_module_callback at
0x7feb7770cca0>
Traceback (most recent call last):
  File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 400, in match_module_callback
    self._make_module_from_path(filepath)
  File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 515, in _make_module_from_path
    module = module_class(filepath, prefix, user_api, internal_api)
  File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 606, in __init__
    self.version = self.get_version()
  File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 646, in get_version
    config = get_config().split()
AttributeError: 'NoneType' object has no attribute 'split'
Exception ignored on calling ctypes callback function: <function _ThreadPoolInfo
._find_modules_with_dl_iterate_phdr.<locals>.match_module_callback at
0x7feb7770cca0>
Traceback (most recent call last):
  File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 400, in match_module_callback
    self._make_module_from_path(filepath)
  File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 515, in _make_module_from_path
    module = module_class(filepath, prefix, user_api, internal_api)
  File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 606, in __init__
    self.version = self.get_version()
  File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 646, in get_version
    config = get_config().split()
AttributeError: 'NoneType' object has no attribute 'split'
Exception ignored on calling ctypes callback function: <function _ThreadPoolInfo
._find_modules_with_dl_iterate_phdr.<locals>.match_module_callback at
0x7feb7770cca0>
Traceback (most recent call last):
  File "/home/musiliandrew/anaconda3/lib/python3.10/site-

```



```

packages/threadpoolctl.py", line 400, in match_module_callback
    self._make_module_from_path(filepath)
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 515, in _make_module_from_path
    module = module_class(filepath, prefix, user_api, internal_api)
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 606, in __init__
    self.version = self.get_version()
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 646, in get_version
    config = get_config().split()
AttributeError: 'NoneType' object has no attribute 'split'
Exception ignored on calling ctypes callback function: <function _ThreadPoolInfo
._find_modules_with_dl_iterate_phdr.<locals>.match_module_callback at
0x7feb7770cca0>
Traceback (most recent call last):
  File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 400, in match_module_callback
    self._make_module_from_path(filepath)
  File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 515, in _make_module_from_path
    module = module_class(filepath, prefix, user_api, internal_api)
  File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 606, in __init__
    self.version = self.get_version()
  File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 646, in get_version
    config = get_config().split()
AttributeError: 'NoneType' object has no attribute 'split'
Exception ignored on calling ctypes callback function: <function _ThreadPoolInfo
._find_modules_with_dl_iterate_phdr.<locals>.match_module_callback at
0x7feb9d86d7e0>
Traceback (most recent call last):
  File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 400, in match_module_callback
    self._make_module_from_path(filepath)
  File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 515, in _make_module_from_path
    module = module_class(filepath, prefix, user_api, internal_api)
  File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 606, in __init__
    self.version = self.get_version()
  File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 646, in get_version
    config = get_config().split()
AttributeError: 'NoneType' object has no attribute 'split'
Exception ignored on calling ctypes callback function: <function _ThreadPoolInfo
._find_modules_with_dl_iterate_phdr.<locals>.match_module_callback at

```

0x7feb7770cca0>

Traceback (most recent call last):

```
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 400, in match_module_callback
    self._make_module_from_path(filepath)
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 515, in _make_module_from_path
    module = module_class(filepath, prefix, user_api, internal_api)
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 606, in __init__
    self.version = self.get_version()
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 646, in get_version
    config = get_config().split()
```

AttributeError: 'NoneType' object has no attribute 'split'

Exception ignored on calling ctypes callback function: <function _ThreadpoolInfo
._find_modules_with_dl_iterate_phdr.<locals>.match_module_callback at

0x7feb7770cca0>

Traceback (most recent call last):

```
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 400, in match_module_callback
    self._make_module_from_path(filepath)
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 515, in _make_module_from_path
    module = module_class(filepath, prefix, user_api, internal_api)
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 606, in __init__
    self.version = self.get_version()
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 646, in get_version
    config = get_config().split()
```

AttributeError: 'NoneType' object has no attribute 'split'

Exception ignored on calling ctypes callback function: <function _ThreadpoolInfo
._find_modules_with_dl_iterate_phdr.<locals>.match_module_callback at

0x7feb9d86d7e0>

Traceback (most recent call last):

```
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 400, in match_module_callback
    self._make_module_from_path(filepath)
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 515, in _make_module_from_path
    module = module_class(filepath, prefix, user_api, internal_api)
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 606, in __init__
    self.version = self.get_version()
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 646, in get_version
    config = get_config().split()
```

```
AttributeError: 'NoneType' object has no attribute 'split'
Exception ignored on calling ctypes callback function: <function _ThreadpoolInfo
._find_modules_with_dl_iterate_phdr.<locals>.match_module_callback at
0x7feb7770cca0>
```

```
Traceback (most recent call last):
```

```
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 400, in match_module_callback
    self._make_module_from_path(filepath)
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 515, in _make_module_from_path
    module = module_class(filepath, prefix, user_api, internal_api)
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 606, in __init__
    self.version = self.get_version()
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 646, in get_version
    config = get_config().split()
```

```
AttributeError: 'NoneType' object has no attribute 'split'
Exception ignored on calling ctypes callback function: <function _ThreadpoolInfo
._find_modules_with_dl_iterate_phdr.<locals>.match_module_callback at
0x7feb7770cca0>
```


```
Traceback (most recent call last):
```

```
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 400, in match_module_callback
    self._make_module_from_path(filepath)
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 515, in _make_module_from_path
    module = module_class(filepath, prefix, user_api, internal_api)
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 606, in __init__
    self.version = self.get_version()
File "/home/musiliandrew/anaconda3/lib/python3.10/site-
packages/threadpoolctl.py", line 646, in get_version
    config = get_config().split()
```

```
AttributeError: 'NoneType' object has no attribute 'split'
```

```
Cluster Centers:
```

```
[[ -0.51458263  0.68771988  0.78079648]
 [ 0.62291581 -0.83250301 -0.94517468]]
```



```
[44]: from scipy.spatial.distance import cdist
from sklearn.metrics import pairwise_distances

def dunn_index(X, labels):
    unique_clusters = np.unique(labels)
    intra_cluster_distances = []
    inter_cluster_distances = []

    # Compute intra-cluster distances (maximum distance within each cluster)
    for cluster in unique_clusters:
        points = X[labels == cluster]
        if len(points) > 1:
            intra_cluster_distances.append(np.max(pairwise_distances(points)))

    # Compute inter-cluster distances (minimum distance between clusters)
    for i in range(len(unique_clusters)):
        for j in range(i + 1, len(unique_clusters)):
            points_i = X[labels == unique_clusters[i]]
            points_j = X[labels == unique_clusters[j]]
            inter_cluster_distances.append(np.min(cdist(points_i, points_j)))
```

```

        return np.min(inter_cluster_distances) / np.max(intra_cluster_distances)

# Compute Dunn Index
dunn_value = dunn_index(X_scaled, kmeans.labels_)
print("Dunn Index:", dunn_value)

```

Dunn Index: 0.2315306889122367

[]:

```

[46]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
from factor_analyzer import FactorAnalyzer

# Load dataset
df = pd.read_csv(get_file_path("2025_FactorAnalysis_Dataset.csv")) # Replace
    ↪ with actual file

# Step 1: Check for missing values and handle them
df.dropna(inplace=True) # Dropping missing values for simplicity

# Step 2: Standardize the data (PCA is affected by scale)
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
df_scaled = scaler.fit_transform(df)

# Step 3: Determine the number of factors using PCA
pca = PCA()
pca.fit(df_scaled)
explained_variance = pca.explained_variance_ratio_

# Plot Scree plot
plt.figure(figsize=(8,5))
plt.plot(range(1, len(explained_variance) + 1), explained_variance.cumsum(),
    ↪ marker='o', linestyle='--')
plt.xlabel("Number of Factors")
plt.ylabel("Cumulative Variance Explained")
plt.title("Scree Plot for Factor Selection")
plt.grid()
plt.show()

# Select factors where cumulative variance reaches ~70% (common threshold)
n_factors = np.argmax(explained_variance.cumsum() >= 0.70) + 1
print(f"Selected Number of Factors: {n_factors}")

```

```

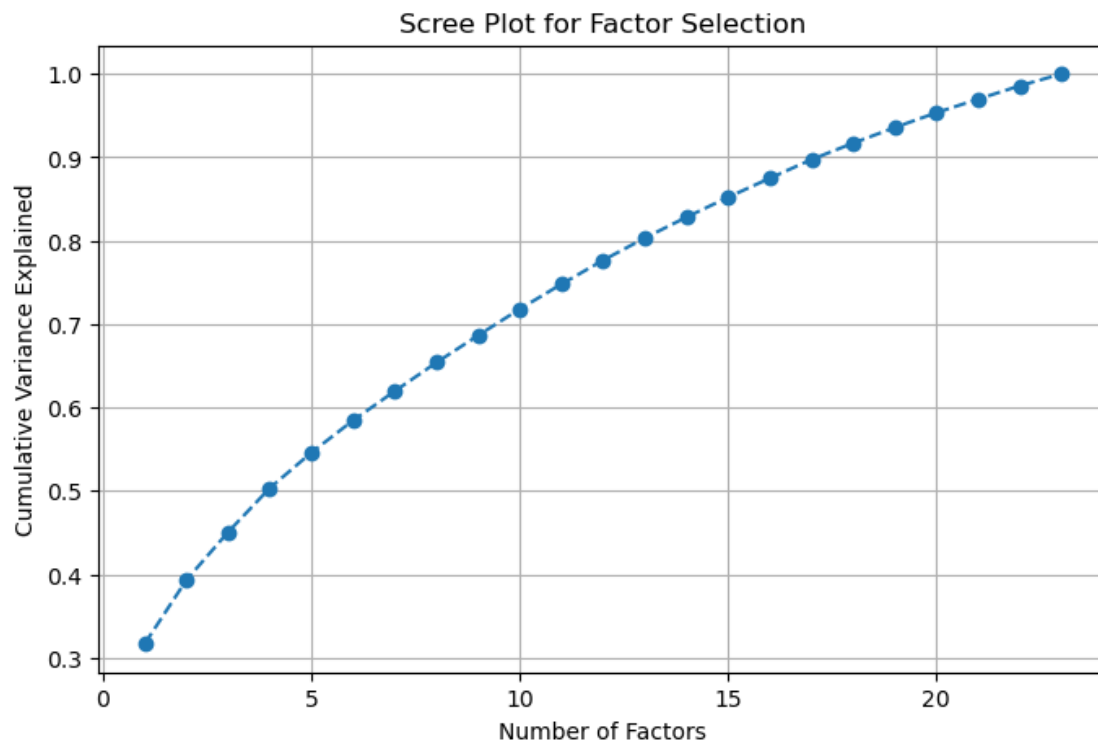
# Step 4: Compute the Cumulative Variance for Selected Factors
cumulative_variance = explained_variance[:n_factors].sum()
print(f"Cumulative Variance Explained by Selected Factors: {cumulative_variance:
↪.2%}")

# Step 5: Perform PCA with Varimax Rotation
fa = FactorAnalyzer(n_factors, rotation="varimax")
fa.fit(df_scaled)
factor_loadings = fa.loadings_

# Step 6: Interpretation - Print Factor Loadings
factor_loadings_df = pd.DataFrame(factor_loadings, index=df.columns,
↪columns=[f'Factor {i+1}' for i in range(n_factors)])
print("\nFactor Loadings after Varimax Rotation:")
print(factor_loadings_df)

# Identify variables highly associated with each factor
for i in range(n_factors):
    associated_vars = factor_loadings_df.iloc[:, i].abs().
↪sort_values(ascending=False).index[:3] # Top 3 variables per factor
    print(f"\nTop Variables Associated with Factor {i+1}:
↪{list(associated_vars)}")

```



Selected Number of Factors: 10

Cumulative Variance Explained by Selected Factors: 71.81%

Factor Loadings after Varimax Rotation:

	Factor 1	Factor 2	Factor 3	Factor 4	Factor 5	Factor 6	Factor 7	\
q01	0.177639	0.581011	0.244872	-0.071363	0.096373	0.021116	-0.010911	
q02	-0.047620	-0.079967	-0.011805	0.459670	-0.149186	-0.050072	0.065653	
q03	-0.178200	-0.327755	-0.185804	0.494354	-0.213854	-0.066752	-0.010930	
q04	0.236886	0.557873	0.229563	-0.103305	0.138406	0.093506	0.009499	
q05	0.232120	0.521850	0.160935	-0.085474	0.121079	0.058713	-0.050905	
q06	0.754655	0.084064	0.129639	-0.057663	0.006659	0.099887	-0.054098	
q07	0.575663	0.266667	0.179769	-0.144333	0.163788	0.107179	-0.034660	
q08	0.126021	0.209417	0.764718	0.002217	0.080449	0.062487	-0.027782	
q09	-0.089330	-0.062031	0.051621	0.620680	-0.032541	-0.057504	0.040708	
q10	0.348726	0.208467	0.103830	-0.127250	-0.019253	0.129090	-0.033906	
q11	0.256313	0.186981	0.721658	-0.165682	0.130074	0.082082	-0.034295	
q12	0.380139	0.379269	0.139529	-0.200534	0.227174	0.086131	-0.015991	
q13	0.575100	0.223219	0.252366	-0.132139	0.102761	0.051763	-0.058449	
q14	0.453716	0.279604	0.171203	-0.147544	0.164795	0.127506	-0.044892	
q15	0.286246	0.169792	0.201814	-0.161065	0.090251	0.895901	-0.039504	
q16	0.207664	0.527495	0.200616	-0.199580	0.144719	0.214604	-0.024059	
q17	0.238965	0.253096	0.642162	-0.029374	0.095279	0.109221	-0.039983	
q18	0.616617	0.242376	0.181302	-0.160264	0.152936	0.038816	-0.019938	
q19	-0.160046	-0.121016	-0.087396	0.366953	-0.183599	-0.044411	0.092367	
q20	0.050055	0.138548	0.114359	-0.213396	0.671744	0.050285	-0.001281	
q21	0.299289	0.304111	0.190421	-0.170676	0.517605	0.045351	-0.011784	
q22	-0.095317	-0.052774	-0.046622	0.332056	-0.020735	-0.039312	0.931972	
q23	-0.044307	0.025509	-0.077634	0.264774	0.019961	0.004137	0.150833	
	Factor 8	Factor 9	Factor 10					
q01	0.155748	0.060399	-0.063942					
q02	0.034674	-0.074374	-0.004886					
q03	-0.054201	-0.105854	-0.166899					
q04	-0.068970	0.077279	0.130515					
q05	-0.017671	-0.035600	-0.003033					
q06	-0.047245	-0.115861	0.007996					
q07	0.006246	-0.041872	0.307867					
q08	0.020844	-0.003190	0.057902					
q09	-0.001796	0.013057	0.116626					
q10	0.057729	-0.054041	-0.114570					
q11	0.002661	0.011878	-0.101274					
q12	0.004851	0.312412	0.055931					
q13	0.034037	0.288048	-0.123808					
q14	0.127789	0.167757	0.078584					
q15	0.064579	0.029242	0.017105					
q16	0.487123	0.018530	0.005668					

q17	0.094310	0.081344	0.095827
q18	0.127589	0.176517	0.063069
q19	-0.078742	-0.039607	-0.070734
q20	0.013416	0.015427	-0.048586
q21	0.077138	0.070395	0.207623
q22	0.003813	-0.025029	0.012580
q23	-0.054655	0.052620	-0.056677

Top Variables Associated with Factor 1: ['q06', 'q18', 'q07']

Top Variables Associated with Factor 2: ['q01', 'q04', 'q16']

Top Variables Associated with Factor 3: ['q08', 'q11', 'q17']

Top Variables Associated with Factor 4: ['q09', 'q03', 'q02']

Top Variables Associated with Factor 5: ['q20', 'q21', 'q12']

Top Variables Associated with Factor 6: ['q15', 'q16', 'q10']

Top Variables Associated with Factor 7: ['q22', 'q23', 'q19']

Top Variables Associated with Factor 8: ['q16', 'q01', 'q14']

Top Variables Associated with Factor 9: ['q12', 'q13', 'q18']

Top Variables Associated with Factor 10: ['q07', 'q21', 'q03']